

## 1. Environment

OS: windows 10

VS version: VS 2017

CUDA version: 9.2

GPU: GTX 1050 Ti

Notice: due to the computer is in Chinese system, the outputs in terminal might contains some Chinese, so the corresponding translation is attached at some outputs screenshots.

## 2. Execution steps of the program

- a) Use VS 2017 to open the project
- b) Use ctrl+f7 to compile the .cu files in the folder
- c) Press ctrl+f5 to run the program
- d) Wait for the result (notice that the program last for 6 min in my own computer)

## 3. Design of the program

- a) Inverted page table structure

Since the memory size given is limited in this project, the program decided to implement the LRU list using double linked list.

- i. For each element in the page table, they are seen as a node of double linked list, and their previous and next node information is store in the same row of that element in the table, that is, for the  $i^{\text{th}}$  element, its previous and next nodes' information in “vm->invert\_page\_table[i + vm->PAGE\_ENTRIES \* 2]” and “vm->invert\_page\_table[i + vm->PAGE\_ENTRIES \* 3]” respectively.
- ii. Also, to avoid using extra memory, the program stores the head and tail information of the list in the first two element of the table, since the size of the inverted page table is 1024, which means they will only occupy the lower bits of the elements and will not affect the valid-invalid bit.
- iii. For valid-invalid bit, 0 in the MSB indicate the bit is valid, and 1 means

invalid. For LRU node, 0x80000000 in the previous or next nodes' information means the LRU node is not inserted to the LRU list. For head and tail hiding in the first two elements in the list, 0x500 in head or tail means the list is empty.

iv. the codes of initialization of inverted page table:

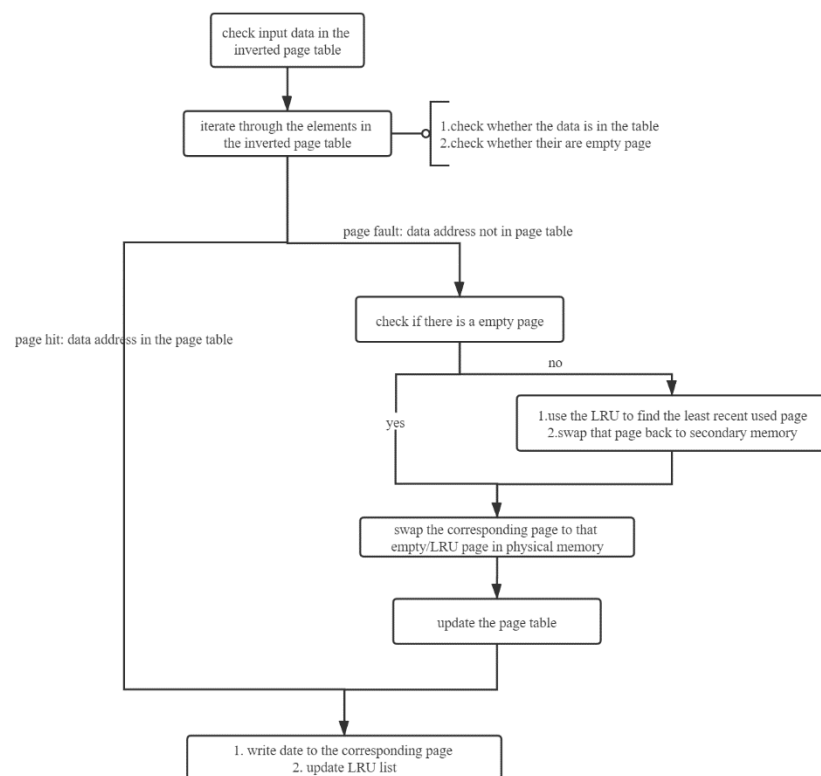
```
device__ void init_invert_page_table(VirtualMemory *vm) {
    for (int i = 0; i < vm->PAGE_ENTRIES; i++) {
        vm->invert_page_table[i] = 0x80000000; // invalid := MSB is 1
        vm->invert_page_table[i + vm->PAGE_ENTRIES] = i;

        // initialize the LRU double linked list
        vm->invert_page_table[i + vm->PAGE_ENTRIES * 2] = 0x80000000;
        vm->invert_page_table[i + vm->PAGE_ENTRIES * 3] = 0x80000000;

        //hide write and tail in the first two elements of invert_page_table
        vm->invert_page_table[0] |= 0x00000500;
        vm->invert_page_table[1] |= 0x00000500;
    }
}
```

b) `vm_write(VirtualMemory *vm, u32 addr, uchar value)`

i. flow chart of the function:



ii. basic logic of the function

1) The function will first go through the invert page table to check whether the data exist in the table and also check whether there exists empty page at the same time.

a) Page hit: If the corresponding page is found, then the iteration will terminate and write data to that address directly.

b) Page fault: in this case, the iteration will go through all the elements to find empty page.

i. If there is an empty page: the program would directly swap the corresponding page in the secondary storage to the frame in physical memory.

ii. Otherwise: it will first use LRU to find the recently least used page and swap that page back to secondary memory. Notice that, to save time, the program directly use that LRU least used page to store the new page, which will directly skip the setting valid bit to invalid and iterating through the table to find an new empty page. After that, the program would swap the corresponding frame in the physical memory.

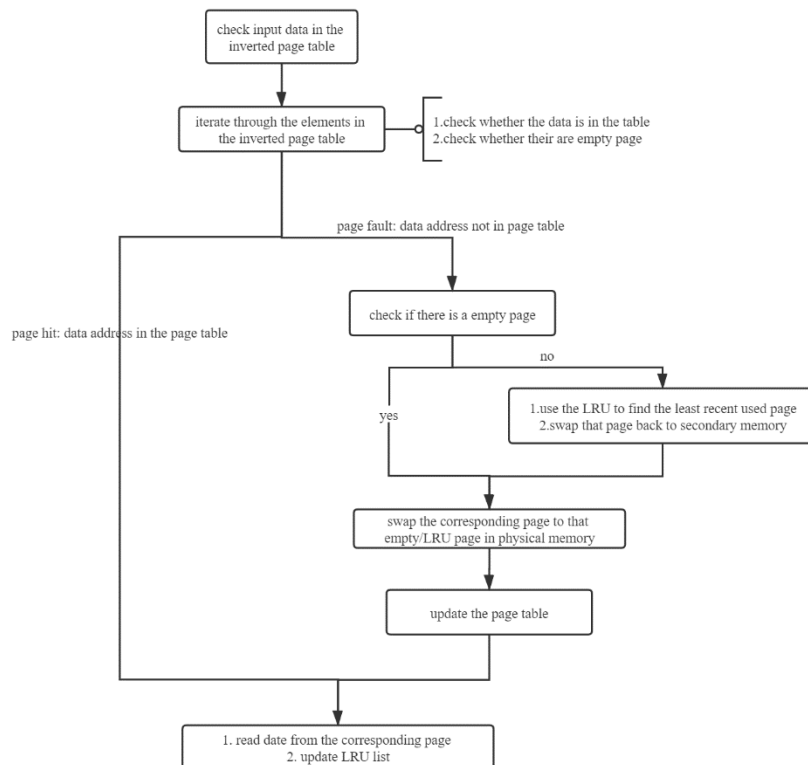
Then the function can directly write the data into the new page.

And update the address of that page in the page table.

2) After writing back the data, the function would set the LRU node of the corresponding element to the front of the LRU list. And also update the head and tail information in the page table.

c) `vm_read(VirtualMemory *vm, u32 addr)`

i. flow chart of the function



## ii. basic logic of the function

The logic is quite similar to the `vm_write()` function, except that the function is reading data instead of writing data.

d) `vm_snapshot(VirtualMemory *vm, uchar *results, int offset, int input_size)`

This function will use `vm_read()` to transfer the data of the address from `offset` to `input_size` to the `results` buffer.

## 4. Page fault of the program:

1) We can modify the `user_program` to show the page fault of each stages clearly.

We can print the page fault of different stages of the function:

```

device__ void user_program(VirtualMemory *vm, uchar *input, uchar *results,
                           int input_size) {

    for (int i = 0; i < input_size; i++) {
        vm_write(vm, i, input[i]);
    }
    printf("finishing writing: page fault -> %d\n", (int) *(vm->pagefault_num_ptr));

    for (int i = input_size - 1; i >= input_size - 32769; i--) {
        int value = vm_read(vm, i);
    }
    printf("finishing reading: page fault -> %d\n", (int) *(vm->pagefault_num_ptr));

    vm_snapshot(vm, results, 0, input_size);
}

```

## 2) outputs:

```

Microsoft Visual Studio 调试控制台
input size: 131072
finishing writing: page fault -> 4096
finishing reading: page fault -> 4097
pagefault number is 8193
D:\CSC3150\assignments\assignment3\A3_template\source\CSC3150_A3\x64\Debug\CSC3150_A3.exe (进程 14616) 已退出, 返回代码为 0。
按任意键关闭此窗口...

```

(the last line of Chinese in the terminal means the programs execute successfully)

This result shows that the final page fault numbers can be decomposed into  $4096+1+4096$ .

- a. In the first write iteration, the program writes the data of address from 0 to 131071 consecutively. Every time when writing the page that is multiple of 32, (like 0, 32, ...), the function would generate a page fault. After that, the next 31 data would not generate page fault, since they are in the same page of that data of address in the multiple of 32. Totally, there are  $131072/32 = 4096$  of that kind of data (whose address of is multiple of 32), which will generate 4096 page faults.
- b. After writing, the program would read 32769 pieces of data from address 131071 consecutively in the decreasing order. Since the first 1024 pages of the last 32768 pieces of data which means they will not generate any page fault. However, for the 32769<sup>th</sup> data, its page does not in the table, which will generate 1 page fault to bring that page in physical memory.
- c. The `vm_snapshot` function will transfer the date of all 131072 data from address 0 consecutively in the increasing order. Notice that since it is transfer data from address 0, the previous 1024 pages in the page table won't get any page hit since the pages are set for last 1024 pages of data.

In this case, the situation is quite similar to the vm\_write process. Then it will also generate 4096 page faults.

## 5. Outputs (with the original user\_program( ))

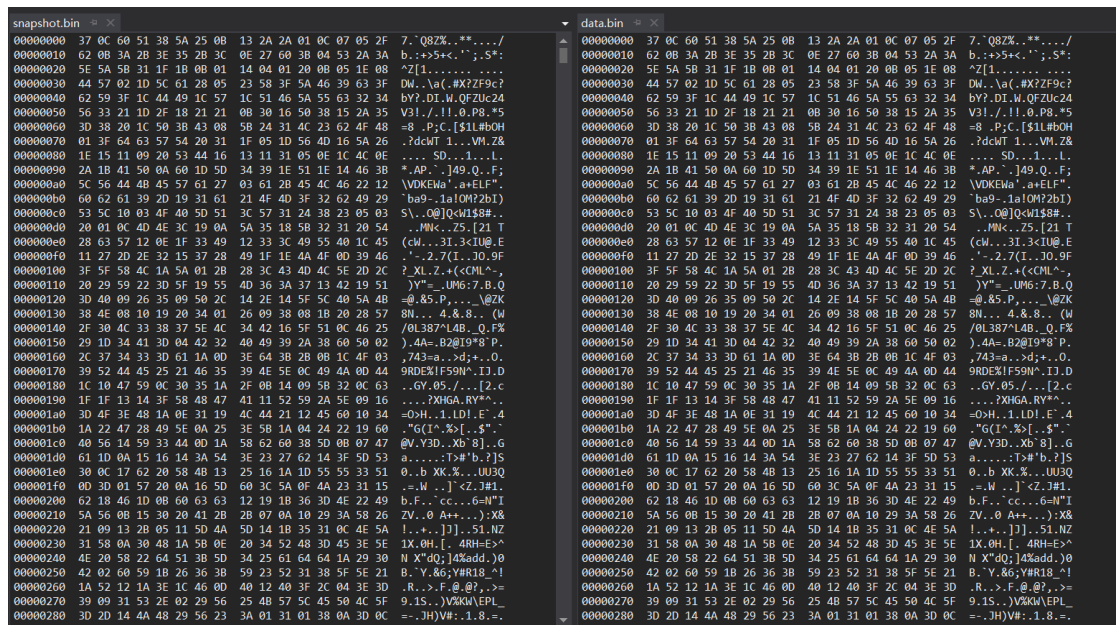
### a. terminal:

```
Microsoft Visual Studio 调试控制台
input size: 131072
pagefault number is 8193

D:\CSC3150\assignments\assignment3\A3_template\source\CSC3150_A3\Debug\CSC3150_A3.exe (进程 26884) 已退出, 返回代码为: 0。
按任意键关闭此窗口...
```

(the last line of Chinese in the terminal means the programs execute successfully)

### b. snapshot.bin



file compare result:

```
D:\CSC3150\assignments\assignment3\A3_template\source\CSC3150_A3>fc data.bin snapshot.bin
正在比较文件 data.bin 和 SNAPSHOT.BIN
FC: 找不到差异

D:\CSC3150\assignments\assignment3\A3_template\source\CSC3150_A3>
```

(the result means there are no difference between two files)

## 6. Problem met in the project

When I implemented the first version of the project, it last about 15 min to output the final result. And I found that the main reason for that is I need to go through

all the elements in page tables at every function call to search for the head and tail of the LRU linked list. Then I decided to hide them in the lower bit in the first two elements of the page table. In this way, it save a lot of time to search for the head and tail of the LRU list and reduce the running time a lot.

#### 7. What I learn from the project

Firstly, I learnt some basic knowledge about cuda programming like declaring variables and functions in different positions using `_device_` and `_host_`. Also, I learnt about the basic structure and mechanism of the virtual memory, and also the detail implementation of simple invert page table. I also learnt about the LRU mechanism and know how to implement it using double linked list.

#### 8. Bonus

The main structure of virtual memory of the bonus part is same as the project, so we can do some modifications to finish the bonus part.

- a. First, we need to launch 4 cuda threads to run the program, which means that in the main.cu function, the kernel part is needed to be changed like this:

```
mykernel<<<1, 4, INVERT_PAGE_TABLE_SIZE>>>(input_size);
```

- b. After testing, I found that for any function call for `vm_read`, `vm_write` and `vm_snapshot`, the program would use 4 threads to perform the exact same function call. Due to this feature, we can use their thread ID to control some threads when their execution. Since the user program always read or write data in consecutive order, in this case, to maintain the non-preemptive order, we can just require each threads to perform the read or write data of corresponding address, that is, we always allow thread 0, 1, 2, 3 to perform the read/write function for data of address that satisfy  $\text{addr} \% 4 = 0, 1, 2 \text{ and } 3$  respectively. Also, since for each call of a function, only one thread will do the corresponding operations, which means there will be no threads racing for the same resource of memory.

- c. For `vm_read` and `vm_write` function, we can just add the following codes at the beginning to control the threads:

a) `vm_read`

```
_device__ uchar vm_read(VirtualMemory *vm, u32 addr) {
    /* Complete vm_read function to read single element from data buffer */
    //check whether the address is in the invert page table

    __syncthreads();
    if (addr % 4 != ((int)threadIdx.x)) return;
    printf("[thread %d] reading %d\n", threadIdx.x, addr);
}
```

b) `vm_write`

```
_device__ void vm_write(VirtualMemory *vm, u32 addr, uchar value) {
    /* Complete vm_write function to write value into data buffer */

    __syncthreads();
    if (addr % 4 != ((int)threadIdx.x)) return;
    printf("[thread %d] writing %d\n", threadIdx.x, addr);
}
```

- d. For `vm_snapshot`, since the program would use 4 threads to execute it, which means we can just reduce the iteration in it to `1/input_size` and assign different threads with different task.

```
_device__ void vm_snapshot(VirtualMemory *vm, uchar *results, int offset,
int input_size) {
    /* Complete snapshot function together with vm_read to load elements from data
    * to result buffer */
    for (int i = offset; i < input_size / 4; i++) {
        results[i * 4 + (int)threadIdx.x] = vm_read(vm, i * 4 + (int)threadIdx.x);
    }
}
```

e. Outputs:

a) Terminal:

i. Start execution



选择D:\CSC3150\assignments\assignment3\A3\_template\source\CSC3150\_A3\_Bonus

```
input size: 131072
[thread 0] writing 0
[thread 1] writing 1
[thread 2] writing 2
[thread 3] writing 3
[thread 0] writing 4
[thread 1] writing 5
[thread 2] writing 6
[thread 3] writing 7
[thread 0] writing 8
[thread 1] writing 9
[thread 2] writing 10
[thread 3] writing 11
[thread 0] writing 12
```

ii. During execution

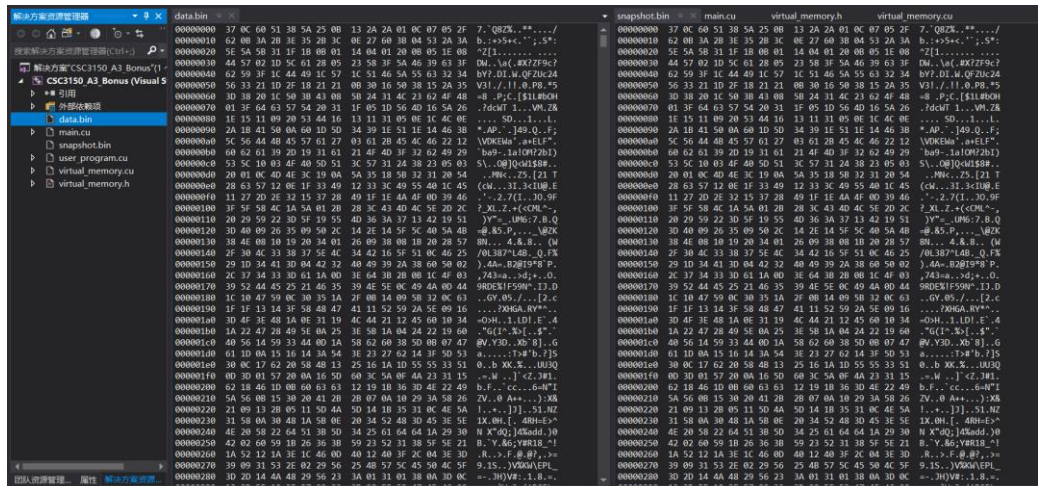
```
[thread 0] writing 8
[thread 1] writing 9
[thread 2] writing 10
[thread 3] writing 11
[thread 0] writing 12
[thread 1] writing 13
[thread 2] writing 14
[thread 3] writing 15
[thread 0] writing 16
[thread 1] writing 17
[thread 2] writing 18
[thread 3] writing 19
```

iii. Finishing execution

```
[thread 0] reading 131064
[thread 1] reading 131065
[thread 2] reading 131066
[thread 3] reading 131067
[thread 0] reading 131068
[thread 1] reading 131069
[thread 2] reading 131070
[thread 3] reading 131071
pagefault number is 8193
```

b) Snapshot.bin:

i. Output file:



ii. File compare:

```
D:\CSC3150\assignments\assignment3\A3_template\source\CSC3150_A3_Bonus\CSC3150_A3_Bonus\fc data.bin snapshot.bin
正在比较文件: data.bin 和 SNAPSHOT.BIN
FC: 找不到差异
```

(the result means there are no difference between two files)